Python Programming for Scientists

Alexander Eberspächer
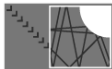
October 12th 2011

OTTO VON GUERICKE
UNIVERSITÄT
MAGDEBURG

NAT

FAKULTÄT FÜR
NATURWISSENSCHAFTEN

FOR760: Scattering Systems with Complex Dynamics

Outline

1 **Introduction**

2 **Basics**

3 **Python Modules for Science**

4 **Faster Python and Glueing**

5 **Summary**

Outline

1 Introduction

2 Basics

3 Python Modules for Science

4 Faster Python and Glueing

5 Summary

Who uses...

- We all use computers to generate or process data

Question to the audience: who uses...

- C/C++?
- Fortran?
- Ada?
- Java?
- Matlab/Octave?

- IDL?
- Perl?
- Ruby?
- **Python**?

What is Python?

Python is/has...

- a scripting language
- general purpose
- interpreted
- easy to learn
- clean syntax

- multi-paradigm
- open-source
- available for all major platforms
- great community

The best of all:

Python comes...

... with **batteries included**!

Libraries available for...

daily IT needs...

- networks
- OS interaction
- temporary files
- zip files
- ...

science!

- efficient array operations (*NumPy*)
- general numerical algorithms (*SciPy*)
- 2D visualization (*matplotlib*)
- 3D visualization (Mayavi)
- special problems (e.g. finite elements with FEniCS, quantum optics with QuTiP)
- symbolic math (SageMath, sympy)
- ...

Outline

Scientific *Hello, World!*

```python
import sys
from math import sin, pi

def sincSquare(x):
    """Return sinc(x)^2.
    """
    if(x <> 0.0):
        return (sin(pi*x)/(pi*x))**2
    else:
        return 1.0

x = sys.argv[1]
y = sincSquare(float(x))
print("sinc(%s)^2 = %s"%(x, y))
```

cmp. H.-P. Langtangen,
"Python Scripting for
Computational Science"

run with:

python HelloWorld.py 0.0

Control structures

```python
# if statements:
if(divisor == 0):
    ...
elif(divisor > 1E20):
    ...
else:
    ...

# loops:
for i range(10): # i = 0, 1, ..., 9
    print("i = %s"%i)

# while loops:
while(True):
    ...
```

Functions

```python
# functions:
def f(x, a=1.0, b=2.0):
    """Return a/x and a/x^b.
    """

    return a/x, a/x**b

# somewhere else:
a = 5
y1, y2 = f(x, 5.0)
y3, y4 = f(2, b=3.0)
```

Data types

```python
a = 2 # integer
b = 2.0 # float
c = "3.0" # string
d = [1, 2, "three"] # list
e = "1"
print(a*b) # valid, upcasting
print(a*c) # valid, but probably not desired: '3.03.0'
print(b*c) # invalid
print(d[1]) # prints 2
for item in d: # lists are "iterable"
    print(item)
for character in c: # strings are iterable
    print(character) # prints 3\n. \n0
f = e + c # + joins strings: f = '13.0'
g = d + [someObj, "foobar"] # + joins lists
```

Introduction
000

Basics
0000●00

Python Modules for Science
000000000

Faster Python and Glueing
0000000000000

Summary
00

Files

```python
readFile = open("infile", mode="r")
writeFile = open("outfile", mode="w")

for line in readFile: # iterate over file's lines
    xString, yString = line.split() # split the line
    x = float(xString); y = float(yString)
    print("x = %s, y = %s"%(x, y))
    writeFile.write("%s * %s = %s\n"%(x, y, x*y))

readFile.close(); writeFile.close()
```

infile:

```
1.0    2.0
3.0    4.0
```

outfile:

```
1.0 * 2.0 = 2.0
3.0 * 4.0 = 12.0
```

Some Python magic

```python
x, y = y, x # swapping

print(1 > 2 > 3) # prints False

# filtering (there is also reduce(), map())
numbers = range(50)
evenNumbers = filter(lambda x: x % 2 == 0, numbers)
print("All even numbers in [0; 50): %s"%evenNumbers)

# list comprehensions:
squares = [x**2 for x in numbers]

a += 2 # a = a + 2

print("string" in "Long string") # prints True
```

Pitfalls

↯ Common pitfalls:

- slicing: last index is *exclusive*, not *inclusive* as in e.g. Fortran

```
x = [1, 2, 3, 4]
print(x[0:2]) # prints [1, 2], not [1, 2, 3]
```

- What looks like an assignment is actually setting a reference:

```
a = []
b = a
a.append(2)
print(a) # prints [2]
print(b) # prints [2], not []!
```

Outline

1 Introduction

2 Basics

## 3 Python Modules for Science

4 Faster Python and Glueing

5 Summary

The IPython shell

IPython

An interactive shell - may replace MatLab [tm] for interactive work

- Syntax highlighting
- Tab completion
- Inline documentation
- Easy profiling, timing...
- IPython ≥ 0.11: inline plots...

NumPy: Python meets an array data type

NumPy

Fast and convenient array operations

- Lists: + does join, not add!
- NumPy array: basic vector/matrix data type
- Convenience functions (e.g. `linspace()`, `zeros()`, `loadtxt()`...)
- Array slicing
- element-wise operations
- Code using NumPy reads and writes very similar to modern Fortran (slicing, vector valued indices...)

NumPy by examples

```python
import numpy as np

a = np.array([1.0, 2.0, 3.0, 4.0])
b = np.array([4.0, 3.0, 2.0, 1.0])
for item in a: # arrays are iterable
    print(item)
c = a + b # c = [5, 5, 5, 5]
print(a[0:3:2]) # 1.0, 3.0; last element not included!
a[0:3] = b[0:-1]

print(a*b) # prints [4, 6, 6, 4], not the scalar product!
```

SciPy

SciPy

Numerical algorithms using NumPy arrays

Wrappers around well-established libraries

Submodules:

- `linalg`: Linear algebra (lapack)
- `sparse`: sparse matrices
- `fft`: FFT (fftpack)
- `optimize`: Optimization, Zeros (minpack)

- `integration`: Integration (quadpack, odepack)
- `special`: special functions (amos...)
- `signal`: Signal processing

SciPy: an example

```python
import numpy as np
from scipy.optimize import curve_fit
from matplotlib.pyplot import plot, show, legend

x, yExp = np.loadtxt("func.dat", unpack=True)
plot(x, yExp, ls="--", c="blue", lw="1.5", label="Exp.")

def fitFunc(x, a, b, c):
    return a*np.exp(-b*x) + c

pOpt, pCov = curve_fit(fitFunc, x, yExp)
yFit = fitFunc(x, a=pOpt[0], b=pOpt[1], c=pOpt[2])
plot(x, yFit, label="Fit: $a = %s; b = %s; c= %s$"\
     %(pOpt[0], pOpt[1], pOpt[2]), ls="-", lw="1.5", c="r")
legend(); show()
```

SciPy: the example's output

Already used here: *Matplotlib*

Matplotlib

(mostly) 2D plots

Pylab: MatLab alternative for interactive work

Introduction
ooo

Basics
ooooooo

Python Modules for Science
oooooooo●oo

Faster Python and Glueing
ooooooooooooo

Summary
oo

Some Pylab: the logistic map

```python
from matplotlib.pylab import * # some of NumPy, SciPy, MPL

rVals = 2000; startVal = 0.5
throwAway = 300; samples = 800
vals = zeros(samples-throwAway)

for r in linspace(2.5, 4.0, rVals): # iterate r
    x = startVal
    for s in range(samples):
        x = r*x*(1-x) # logistic map
        if(s >= throwAway): vals[s-throwAway] = x
    scatter(r*ones(samples-throwAway), vals, c="k", \
            marker="o", s=0.3, lw=0) # plot

xlabel("$r$"); ylabel("$x$"); title("Log. map"); show();
```

Some Pylab: the logistic map

The last script produces this image:

Introduction	Basics	Python Modules for Science	**Faster Python and Glueing**	Summary
ooo	ooooooo	ooooooooo	oooooooooooo	oo

Outline

1 Introduction

2 Basics

3 Python Modules for Science

4 Faster Python and Glueing

5 Summary

Using Python as glue

Python can wrap different different other programming languages

Cython

compiled, *typed* Python - interface C/C++ code

f2py

Fortran wrapper, included in NumPy

Why do that?

- Python can be *slow*
- Python loops are slow
- calling Python functions is slow

- Wrap external C/Fortran... libraries
- Happily/unfortunately (?) there is legacy code

Problem: $\mathrm{sinc}(x)^2$

```python
import numpy as np
from math import sin, pi

def sincSquare(x):
    """Return the sinc(x) = (sin(x)/x)**2 of the array
    argument x.
    """
    retVal = np.zeros_like(x)
    for i in range(len(x)):
        retVal[i] = (sin(pi*x[i]) / (pi*x[i]))**2

    return retVal
```

10^6 array elements: 1 loops, best of 3: 4.91 s per loop

Problem: $\text{sinc}(x)^2$

```python
import numpy as np
from math import sin, pi

def sincSquare(x):
    """Return the sinc(x) = (sin(x)/x)**2 of the array
    argument x.
    """
    retVal = np.zeros_like(x)
    for i in range(len(x)):
        retVal[i] = (sin(pi*x[i]) / (pi*x[i]))**2

    return retVal
```

10^6 array elements: 1 loops, best of 3: 4.91 s per loop

First attempt: use NumPy array operations

```
import numpy as np

def sincSquareNumPy1(x):

    return (np.sin(np.pi*x[:])/(np.pi*x[:]))**2

def sincSquareNumPy2(x):

    return np.sinc(x[:])**2
```

10^6 array elements: first function: 10 loops, best of 3: 73 ms
per loop, second function: 10 loops, best of 3: 92.9 ms
per loop

First attempt: use NumPy array operations

```
import numpy as np

def sincSquareNumPy1(x):

    return (np.sin(np.pi*x[:])/(np.pi*x[:]))**2

def sincSquareNumPy2(x):

    return np.sinc(x[:])**2
```

10^6 array elements: first function: 10 loops, best of 3: 73 ms
per loop, second function: 10 loops, best of 3: 92.9 ms
per loop

How Cython works

Cython

compiled, possibly typed Python:

.pyx file $\overset{\text{Cython}}{\Longrightarrow}$ .c file $\overset{\text{C compiler}}{\Longrightarrow}$ .so/.dll file

- various levels of typing possible
- C output and Cython's opinion on code speed can easily be inspected (optional .html output)
- interfacing C libraries is easy

$\mathrm{sinc}(x)^2$ - **Cython, Version 1**

```
cdef extern from "math.h":
    double sin(double)
    double pow(double, int)

def sincSquareCython1(x):

    pi = 3.141592653589793238462643
    retVal = np.zeros_like(x)

    for i in range(len(x)):
        retVal[i] = (sin(pi*x[i]) / (pi*x[i]))**2

    return retVal
```

10^6 array elements: 1 loops, best of 3: 4.39 s per loop

$\mathrm{sinc}(x)^2$ - **Cython, Version 1**

```
cdef extern from "math.h":
    double sin(double)
    double pow(double, int)

def sincSquareCython1(x):

    pi = 3.1415926535897932384626433
    retVal = np.zeros_like(x)

    for i in range(len(x)):
        retVal[i] = (sin(pi*x[i]) / (pi*x[i]))**2

    return retVal
```

10^6 array elements: 1 loops, best of 3: 4.39 s per loop

$\mathrm{sinc}(x)^2$ - **Cython, Version 2**

```python
cimport numpy as np # also C-import types

cpdef np.ndarray[double] sincSquareCython2\
    (np.ndarray[double] x):

    cdef int i
    cdef double pi = 3.1415926535897932384626433
    cdef np.ndarray[double] retVal = np.zeros_like(x)

    for i in range(len(x)):
        retVal[i] = pow(sin(pi*x[i]) / (pi*x[i]), 2)
```

10^6 array elements: 10 loops, best of 3: 49.1 ms per loop
That's a **speedup by a factor** ≈ 100!

$\mathrm{sinc}(x)^2$ - **Cython, Version 2**

```
cimport numpy as np # also C-import types

cpdef np.ndarray[double] sincSquareCython2\
    (np.ndarray[double] x):

    cdef int i
    cdef double pi = 3.1415926535897932384626433
    cdef np.ndarray[double] retVal = np.zeros_like(x)

    for i in range(len(x)):
        retVal[i] = pow(sin(pi*x[i]) / (pi*x[i]), 2)
```

10^6 array elements: 10 loops, best of 3: 49.1 ms per loop

That's a **speedup by a factor ≈ 100**!

How f2py works

f2py

wrap Fortran code in Python:

.f/.f90 file $\overset{\text{f2py}}{\Longrightarrow}$ .so/.dll file

- f2py is included in NumPy
- exposes NumPy arrays to Fortran code
- once 'Fortran space' is entered, you run at full Fortran speed

$\text{sinc}(x)^2$ - **f2py, Version 1**

```
subroutine sincsquaref2py1(x, n, outVal)
    implicit none

    double precision, dimension(n), intent(in) :: x
    integer, intent(in) :: n
    double precision, dimension(n), intent(out) :: outVal
    double precision, parameter :: pi = 4.0d0 * atan(1.0d0)

    outVal(:) = (sin(pi*x(:)) / (pi*x(:)))**2

end subroutine sincsquaref2py1
```

10^6 array elements: 10 loops, best of 3: 47.4 ms per loop
Again, a **speedup by a factor of** ≈ 100!

$\mathrm{sinc}(x)^2$ - **f2py, Version 1**

```fortran
subroutine sincsquaref2py1(x, n, outVal)
    implicit none

    double precision, dimension(n), intent(in) :: x
    integer, intent(in) :: n
    double precision, dimension(n), intent(out) :: outVal
    double precision, parameter :: pi = 4.0d0 * atan(1.0d0)

    outVal(:) = (sin(pi*x(:)) / (pi*x(:)))**2

end subroutine sincsquaref2py1
```

10^6 array elements: 10 loops, best of 3: 47.4 ms per loop
Again, a **speedup by a factor of** ≈ 100!

Introduction	Basics	Python Modules for Science	**Faster Python and Glueing**	Summary
ooo	ooooooo	ooooooooo	oooooooo**o**oo**o**ooo	oo

Cheating: $\mathrm{sinc}(x)^2$ - f2py, Version 2 - OpenMP

```fortran
subroutine sincsquaref2py2(x, n, outVal)
    implicit none
    double precision, dimension(n), intent(in) :: x
    integer, intent(in) :: n
    double precision, dimension(n), intent(out) :: outVal
    integer :: i
    double precision, parameter :: pi = 4.0d0 * atan(1.0d0)
    !$OMP PARALLEL DO SHARED(x, outVal)
    do i = 1, n
        outVal(i) = (sin(pi*x(i)) / (pi*x(i)))**2
    end do
    !$OMP END PARALLEL DO
end subroutine sincsquaref2py2
```

10^6 array elements, 2 Threads: 10 loops, best of 3: 33.5 ms

Cheating: $\mathrm{sinc}(x)^2$ - f2py, Version 2 - OpenMP

```
subroutine sincsquaref2py2(x, n, outVal)
    implicit none
    double precision, dimension(n), intent(in) :: x
    integer, intent(in) :: n
    double precision, dimension(n), intent(out) :: outVal
    integer :: i
    double precision, parameter :: pi = 4.0d0 * atan(1.0d0)
    !$OMP PARALLEL DO SHARED(x, outVal)
    do i = 1, n
        outVal(i) = (sin(pi*x(i)) / (pi*x(i)))**2
    end do
    !$OMP END PARALLEL DO
end subroutine sincsquaref2py2
```

10^6 array elements, 2 Threads: 10 loops, best of 3: 33.5 ms

$\text{sinc}(x)^2$ - **Overview**

Benchmark for an Intel i7:

Time for calls to $\left(\frac{\sin(\pi x)}{\pi x}\right)^2$ for 500000 array elements

Techniques for faster Scripts

After you have written a prototype in Python with NumPy and SciPy, check if your code is already fast enough. If not,

- profile your script (IPython's run -p or cProfile module...) to find bottlenecks
- if a large numbers of function calls is the bottleneck, typing and using Cython's cdef/cpdef for C calling conventions speeds your code up at the cost of flexibility
- loops greatly benefit from typing, too
- consider moving heavy computations to Fortran/C completely - f2py and Cython will help you wrapping

Slightly OffTopic: mpi4py

mpi4py

Interface MPI in Python

- speed-up pure Python by parallelization using MPI (OpenMPI, mpich...)
- mpi4py also works with f2py and Cython (?)
- $\rightarrow$ run the steering Python script with `mpirun...`, take care of the communicator there and use it in Fortran, too

Alternatives:

- IPython's parallel computing facilities

Slightly OffTopic: mpi4py

```python
from mpi4py import MPI

MPIroot = 0 # define the root process
MPIcomm = MPI.COMM_WORLD # MPI communicator

MPIrank, MPIsize = MPIcomm.Get_rank(), MPIcomm.Get_size()

...

MPIcomm.Reduce(tempVals, retVal, op=MPI.SUM, root=MPIroot)
```

Outline

Python in teaching

Python/Pylab should be used in teaching because

- it is easy...
- and yet powerful;
- it may be used specialized to numerical computing...
- and also serve students as a general purpose language;
- it is safe;
- and best of all, it is *free*!

Take home message 1

Python is ideal for teaching

Summary

We have...

- introduced basic Python scripting
- shown some basic modules for scientific computing
- demonstrated how to wrap other languages
- learned how to speed Python up

Take home message 2

Python is a very valuable tool for Physicists

Slides, LATEX and Python Sources available at
http://github.com/aeberspaecher